

Neural Networks

4. Multi-layer Perceptron & Back-propagation

Farkaš, Kuzma et al.

Center for Cognitive Science
Department of Applied Informatics
Faculty of Mathematics, Physics and Informatics UKBA

March 13th, 2018

Conventions

Conventions, vol. 1

c – scalar

\mathbf{x} – column vector (“western” algebra)

x_i – i -th element of a vector (scalar)

\mathbf{W} – matrix

\mathbf{w}_i – i -th row of matrix (vector)

$w_{i,j}$ – element in the i -th row and j -th column (scalar)

Conventions, vol. 2

$c\mathbf{x}, c\mathbf{W}$ – scalar/elementwise multiplication – shape is preserved

$\mathbf{W}\mathbf{X}$ – matrix multiplication

$\mathbf{W}\mathbf{x}$ – matrix-vector multiplication, i.e. applying transformation \mathbf{W} to a vector \mathbf{x}

$\mathbf{x}\mathbf{y}$ – mistake

$\mathbf{x}^T\mathbf{y} = \mathbf{x} \cdot \mathbf{y}$ – scalar/dot/inner product – result is a scalar

$\mathbf{x} \times \mathbf{y}$ – vector product – result is a vector

$\mathbf{x}\mathbf{y}^T$ – outer product – result is a matrix


Numpy

Numpy

- `np.zeros`, `np.ones`, `np.eye`
- uniform `np.random.rand`, normal/gaussian `np.random.randn`
- `+`, `-`, `*`, `np.sin(...)`, `np.exp(...)` all work elementwise
- `+=`, `-=`, `*=` also work
- `np.dot(a,b) = a.dot(b) = a @ b`
- `np.transpose(a) = a.transpose() = a.T`
- transpose on a vector does nothing!
 - use `np.inner` or `np.outer`
 - alternatively, `np.atleast_2d(x).T` returns a matrix posing as a row vector
- `np.linalg.norm`
- `np.min/max/mean/std(x, axis=0)`


Multi-layer Perceptron

Multi-layer Perceptron



- virtual input for the bias term: $x_{n+1} = h_{m+1} = 1$


Activation functions



- logistic sigmoid:
 - $\text{logsig}(x) = \frac{1}{1+e^{-x}}$
 - $\text{logsig}'(x) = \text{logsig}(x)(1 - \text{logsig}(x))$
- hyperbolic tangent
- rectified linear units (ReLU)
- linear (makes sense only on output)
- ...

Back-propagation

Back-propagation



- $g_i^{out} = (d_i - y_i) f'_{out}(b_i)$
- $g_k^{hid} = \left(\sum_i w_{i,k}^{out} g_i^{out} \right) f'_{hid}(a_k)$
- $\Delta W^{out} = \mathbf{g}^{out} \mathbf{h}^T \quad \mathbf{W}^{out}(t+1) = \mathbf{W}^{out}(t) + \alpha \Delta \mathbf{W}^{out}(t)$
- $\Delta W^{hid} = \mathbf{g}^{hid} \mathbf{x}^T \quad \mathbf{W}^{hid}(t+1) = \mathbf{W}^{hid}(t) + \alpha \Delta \mathbf{W}^{hid}(t)$

Algorithm

Initialization:

1. choose model parameters (# of hidden neurons)
2. choose training parameters (learning rate, #epochs)
3. generate random initial weights

Training:

- until stopping criterion (accuracy / #epochs / time...):
 - with each training sample (\mathbf{x}, \mathbf{d}) in random order:
 - forward-pass: compute \mathbf{h}, \mathbf{y}
 - backward-pass: compute $\mathbf{g}^{out}, \mathbf{g}^{hid}$
 - adjust weights $\mathbf{W}^{hid}, \mathbf{W}^{out}$